

```
1 import pandas as pd
2 import seaborn as sns
3 import numpy as np
4 import os
5 import PIL
6 import pickle
7 from PIL import *
8 import cv2
9 import tensorflow as tf
10 import keras
11 from keras.applications import DenseNet121 # 2017 architecture
12 from keras.models import Model, load_model
13 from keras.initializers import glorot_uniform
14 from keras.utils import plot_model
15 from keras.callbacks import ReduceLROnPlateau, EarlyStopping, ModelCheckpoint, LearningRateScheduler
16 from IPython.display import display
17 from keras import *
18 from keras.preprocessing.image import ImageDataGenerator
19 from keras import layers, optimizers
20 from keras.applications.resnet50 import ResNet50
21 from keras.layers import *
22 from keras import backend as K
23 from keras import optimizers
24 import matplotlib.pyplot as plt
25 from sklearn.model_selection import train_test_split
```

```


1 # RESBLOCK :=> input -> Convolution Block -> Identity Block -> Identity Block -> output
2 def res_block(X, filters,stage):
3     # Convolution Block
4     X_copy = X
5
6     f1, f2, f3 = filters
7
8     # Main path
9     X = Conv2D(f1,(1,1), strides = (1,1), name = 'res_'+str(stage)+'_conv_a', kernel_initializer = glorot_uniform(seed = 0))(X)
10    X = MaxPool2D((2,2))(X)
11    X = BatchNormalization(axis = 3, name = 'bn_'+ str(stage) + '_conv_a')(X)
12    X = Activation('relu')(X)
13
14    X = Conv2D(f2,(3,3), strides = (1,1), padding= 'same', name = 'res_' + str(stage) + '_conv_b', kernel_initializer = glorot_uniform(se
15
16    X = BatchNormalization(axis = 3, name = 'bn_'+ str(stage) + '_conv_b')(X)
17    X = Activation('relu')(X)
18
19    X = Conv2D(f3,(1,1),strides = (1,1), name = 'res_' + str(stage) + '_conv_c', kernel_initializer = glorot_uniform(seed = 0))(X)
20    X = BatchNormalization(axis = 3, name = 'bn_'+ str(stage) + '_conv_c')(X)
21
22    # Short path
23
24    X_copy = Conv2D(f3,(1,1), strides =(1,1), name = 'res_'+str(stage) + '_conv_copy', kernel_initializer = glorot_uniform(seed = 0))(X_cc
25    X_copy = MaxPool2D((2,2))(X_copy)
26    X_copy = BatchNormalization(axis = 3, name='bn_'+str(stage)+'_conv_copy')(X_copy)
27
28    # ADD
29    X = Add()([X,X_copy])
30    X = Activation('relu')(X)
31
32    # Identity BLock 1
33    X_copy = X
34
35    # main path
36
37    X = Conv2D(f1,(1,1), strides = (1,1), name = 'res_'+str(stage)+'_identity_1_a', kernel_initializer = glorot_uniform(seed = 0))(X)
38    X = BatchNormalization(axis = 3, name = 'bn_'+ str(stage) + '_identity_1_a')(X)
39    X = Activation('relu')(X)
40
41    X = Conv2D(f2,(3,3), strides = (1,1), padding= 'same', name = 'res_' + str(stage) + '_identity_1_b', kernel_initializer = glorot_unif
42    X = BatchNormalization(axis = 3, name = 'bn_'+ str(stage) + '_identity_1_b')(X)
43    X = Activation('relu')(X)
44
45    X = Conv2D(f3, (1,1),strides = (1,1), name = 'res_' + str(stage) + '_identity_1_c', kernel_initializer = glorot_uniform(seed = 0))(X)
46    X = BatchNormalization(axis = 3, name = 'bn_'+ str(stage) + '_identity_1_c')(X)
47
48    # ADD
49    X = Add()([X,X_copy])
50    X = Activation('relu')(X)
51
52    # identity block 2
53    X_copy = X
54
55    # main path
56    X = Conv2D(f1,(1,1), strides = (1,1), name = 'res_'+str(stage)+'_identity_2_a', kernel_initializer = glorot_uniform(seed = 0))(X)
57    X = BatchNormalization(axis = 3, name = 'bn_'+ str(stage) + '_identity_2_a')(X)
58    X = Activation('relu')(X)
59
60    X = Conv2D(f2,(3,3), strides = (1,1), padding= 'same', name = 'res_' + str(stage) + '_identity_2_b', kernel_initializer = glorot_unif
61    X = BatchNormalization(axis = 3, name = 'bn_'+ str(stage) + '_identity_2_b')(X)
62    X = Activation('relu')(X)
63
64    X = Conv2D(f3, (1,1),strides = (1,1), name = 'res_' + str(stage) + '_identity_2_c', kernel_initializer = glorot_uniform(seed = 0))(X)
65    X = BatchNormalization(axis = 3, name = 'bn_'+ str(stage) + '_identity_2_c')(X)
66    # ADD
67    X = Add()([X,X_copy])
68    X = Activation('relu')(X)
69
70    return X

```

```

1 from google.colab import drive
2 drive.mount('/content/drive')
3 facialexpression_df = pd.read_csv('/content/drive/MyDrive/DL Facial Recognition/icml_face_data.csv')

```

 Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
1 facialexpression_df = facialexpression_df.drop(columns = " Usage")
2
```

```
1 facialexpression_df.head()
```

```

↗
   emotion                                     pixels
0         0  70 80 82 72 58 58 60 63 54 58 60 48 89 115 121...
1         0 151 150 147 155 148 133 111 140 170 174 182 15...
2         2 231 212 156 164 174 138 161 173 182 200 106 38...
3         4  24 32 36 30 32 23 19 20 30 41 21 22 32 34 21 1...
4         6    4 0 0 0 0 0 0 0 0 0 0 3 15 23 28 48 50 58 84...

```

```
1 def string2array(x):
2     return np.array(x.split(' ')).reshape(48,48,1).astype('float32')
```

```
1 Start coding or generate with AI.
```

```
1 def resize(x):
2     img = x.reshape(48,48)
3     return cv2.resize(img, dsize=(96,96), interpolation = cv2.INTER_CUBIC)
```

```
1 facialexpression_df[' pixels'] = facialexpression_df[' pixels'].apply(lambda x : string2array(x))
2 facialexpression_df[' pixels'] = facialexpression_df[' pixels'].apply(lambda x : resize(x))
```

```
1 facialexpression_df.shape
```

```
↗ (35887, 2)
```

```
1 label_to_text = {0: 'anger', 1 : 'disgust', 2 : 'fear', 3 : 'happiness', 4: 'sad', 5: 'surprise', 6: 'neutral'}
```

```

1 emotions = [1]
2 for i in emotions:
3     data = facialexpression_df[facialexpression_df['emotion']==i][:1]
4
5     img = data[' pixels'].item()
6     img = img.reshape(96, 96)
7
8     plt.figure()
9     plt.title(label_to_text[i])
10    plt.imshow(img, cmap = 'gray')
11    plt.axis('off')

```

```
↗
      disgust
```



```
1 facialexpression_df.emotion.value_counts().index
```

```
Index([3, 6, 4, 2, 0, 5, 1], dtype='int64', name='emotion')
```

```
1 facialexpression_df.emotion.value_counts()
```

```
emotion
3    8989
6    6198
4    6077
2    5121
0    4953
5    4002
1     547
Name: count, dtype: int64
```

```
1 class1 = facialexpression_df.loc[facialexpression_df['emotion'] == 1]
2 class1.shape
```

```
(547, 2)
```

```
1 from keras.utils import to_categorical
2
3 X = facialexpression_df['pixels']
4 y = to_categorical(facialexpression_df['emotion'])
5
6 X = np.stack(X, axis = 0)
7 X = X.reshape(35887,96,96,1)
8
9 print(X.shape, y.shape)
```

```
(35887, 96, 96, 1) (35887, 7)
```

```
1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, shuffle = True)
3 X_val, X_test, y_val, y_test = train_test_split(X_test, y_test, test_size = 0.5, shuffle = True)
```

```
1 X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
((32298, 96, 96, 1), (1795, 96, 96, 1), (32298, 7), (1795, 7))
```

```
1 X_train = X_train/255
2 X_val   = X_val /255
3 X_test  = X_test/255
```

```
1 train_datagen = ImageDataGenerator(
2     rotation_range = 15,
3     width_shift_range = 0.1,
4     height_shift_range = 0.1,
5     shear_range = 0.1,
6     zoom_range = 0.1,
7     horizontal_flip = True,
8     fill_mode = "nearest")
```

```

1 input_shape = (96, 96, 1)
2
3 # Input tensor shape
4 X_input = Input(input_shape)
5
6 # Zero-padding
7 X = ZeroPadding2D((3, 3))(X_input)
8
9 # 1 - stage
10 X = Conv2D(64, (7, 7), strides= (2, 2), name = 'conv1', kernel_initializer= glorot_uniform(seed = 0))(X)
11 X = BatchNormalization(axis =3, name = 'bn_conv1')(X)
12 X = Activation('relu')(X)
13 X = MaxPooling2D((3, 3), strides= (2, 2))(X)
14
15 # 2 - stage
16 X = res_block(X, [64, 64, 256], stage= 2)
17
18 # 3 - stage
19 X = res_block(X, [128, 128, 512], stage= 3)
20
21 # 4 - stage
22 X = res_block(X, filter= [256, 256, 1024], stage= 4)
23
24 # Average Pooling
25 X = AveragePooling2D((4, 4), name = 'Averagea_Pooling')(X)
26
27 # Final layer
28 X = Flatten()(X)
29 X = Dense(7, activation = 'softmax', name = 'Dense_final', kernel_initializer= glorot_uniform(seed=0))(X)
30
31 facialmodel = Model(inputs= X_input, outputs = X, name = 'Resnet18')
32
33 facialmodel.summary()

```

Model: "Resnet18"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 96, 96, 1)]	0	[]
zero_padding2d (ZeroPaddin g2D)	(None, 102, 102, 1)	0	['input_1[0][0]']
conv1 (Conv2D)	(None, 48, 48, 64)	3200	['zero_padding2d[0][0]']
bn_conv1 (BatchNormalizati on)	(None, 48, 48, 64)	256	['conv1[0][0]']
activation (Activation)	(None, 48, 48, 64)	0	['bn_conv1[0][0]']
max_pooling2d (MaxPooling2 D)	(None, 23, 23, 64)	0	['activation[0][0]']
res_2_conv_a (Conv2D)	(None, 23, 23, 64)	4160	['max_pooling2d[0][0]']
max_pooling2d_1 (MaxPoolin g2D)	(None, 11, 11, 64)	0	['res_2_conv_a[0][0]']
bn_2_conv_a (BatchNormaliz ation)	(None, 11, 11, 64)	256	['max_pooling2d_1[0][0]']
activation_1 (Activation)	(None, 11, 11, 64)	0	['bn_2_conv_a[0][0]']
res_2_conv_b (Conv2D)	(None, 11, 11, 64)	36928	['activation_1[0][0]']
bn_2_conv_b (BatchNormaliz ation)	(None, 11, 11, 64)	256	['res_2_conv_b[0][0]']
activation_2 (Activation)	(None, 11, 11, 64)	0	['bn_2_conv_b[0][0]']
res_2_conv_copy (Conv2D)	(None, 23, 23, 256)	16640	['max_pooling2d[0][0]']
res_2_conv_c (Conv2D)	(None, 11, 11, 256)	16640	['activation_2[0][0]']
max_pooling2d_2 (MaxPoolin g2D)	(None, 11, 11, 256)	0	['res_2_conv_copy[0][0]']
bn_2_conv_c (BatchNormaliz ation)	(None, 11, 11, 256)	1024	['res_2_conv_c[0][0]']
bn_2_conv_copy (BatchNorma	(None, 11, 11, 256)	1024	['max_pooling2d_2[0][0]']

```

lization)

add (Add)                (None, 11, 11, 256)      0      ['bn_2_conv_c[0][0]',
                                              'bn_2_conv_copy[0][0]']

activation_3 (Activation) (None, 11, 11, 256)      0      ['add[0][0]']

res_2_identity_1_a (Conv2D (None, 11, 11, 64)      16448   ['activation_3[0][0]']
)

bn_2_identity_1_a (BatchNorm (None, 11, 11, 64)      256     ['res_2_identity_1_a[0][0]']

1 facialmodel.compile(optimizer=keras.optimizers.Adam(learning_rate=1e-3), loss = "categorical_crossentropy", metrics = ["accuracy"])

1 # using early stopping to exit training if validation loss is not decreasing even after certain epochs (patience)
2 earllystopping = EarlyStopping(monitor = 'val_loss', mode = 'min', verbose = 1, patience = 20)
3
4 # save the best model with lower validation loss
5 checkpointer = ModelCheckpoint(filepath = "/content/drive/MyDrive/DL Facial Recognition /FacialExpression_weights.hdf5", verbose = 1, sav

1 history = facialmodel.fit(train_datagen.flow(X_train, y_train, batch_size=64),
2     validation_data= (X_val, y_val), steps_per_epoch=len(X_train) // 64,
3     epochs= 100, callbacks=[checkpointer, earllystopping])
504/504 [=====] - ETA: 0s - loss: 0.7507 - accuracy: 0.7022
Epoch 27: val_loss did not improve from 0.98270
504/504 [=====] - 29s 58ms/step - loss: 0.7907 - accuracy: 0.7022 - val_loss: 0.9974 - val_accuracy: 0.6282
Epoch 28/100
504/504 [=====] - ETA: 0s - loss: 0.7845 - accuracy: 0.7062
Epoch 28: val_loss did not improve from 0.98270
504/504 [=====] - 29s 57ms/step - loss: 0.7845 - accuracy: 0.7062 - val_loss: 1.0917 - val_accuracy: 0.5948
Epoch 29/100
504/504 [=====] - ETA: 0s - loss: 0.7756 - accuracy: 0.7090
Epoch 29: val_loss improved from 0.98270 to 0.97295, saving model to /content/drive/MyDrive/DL Facial Recognition /FacialExpression_w
504/504 [=====] - 30s 59ms/step - loss: 0.7756 - accuracy: 0.7090 - val_loss: 0.9729 - val_accuracy: 0.6382
Epoch 30/100
504/504 [=====] - ETA: 0s - loss: 0.7640 - accuracy: 0.7163
Epoch 30: val_loss did not improve from 0.97295
504/504 [=====] - 29s 57ms/step - loss: 0.7640 - accuracy: 0.7163 - val_loss: 1.2254 - val_accuracy: 0.5368
Epoch 31/100
504/504 [=====] - ETA: 0s - loss: 0.7582 - accuracy: 0.7155
Epoch 31: val_loss did not improve from 0.97295
504/504 [=====] - 29s 57ms/step - loss: 0.7582 - accuracy: 0.7155 - val_loss: 1.0067 - val_accuracy: 0.6249
Epoch 32/100
504/504 [=====] - ETA: 0s - loss: 0.7490 - accuracy: 0.7198
Epoch 32: val_loss did not improve from 0.97295
504/504 [=====] - 29s 57ms/step - loss: 0.7490 - accuracy: 0.7198 - val_loss: 1.0215 - val_accuracy: 0.6137
Epoch 33/100
504/504 [=====] - ETA: 0s - loss: 0.7392 - accuracy: 0.7228
Epoch 33: val_loss did not improve from 0.97295
504/504 [=====] - 29s 57ms/step - loss: 0.7392 - accuracy: 0.7228 - val_loss: 0.9828 - val_accuracy: 0.6349
Epoch 34/100
504/504 [=====] - ETA: 0s - loss: 0.7309 - accuracy: 0.7254
Epoch 34: val_loss did not improve from 0.97295
504/504 [=====] - 29s 58ms/step - loss: 0.7309 - accuracy: 0.7254 - val_loss: 1.1056 - val_accuracy: 0.6193
Epoch 35/100
504/504 [=====] - ETA: 0s - loss: 0.7267 - accuracy: 0.7284
Epoch 35: val_loss did not improve from 0.97295
504/504 [=====] - 29s 57ms/step - loss: 0.7267 - accuracy: 0.7284 - val_loss: 1.1731 - val_accuracy: 0.5680
Epoch 36/100
504/504 [=====] - ETA: 0s - loss: 0.7145 - accuracy: 0.7351
Epoch 36: val_loss did not improve from 0.97295
504/504 [=====] - 29s 57ms/step - loss: 0.7145 - accuracy: 0.7351 - val_loss: 1.6655 - val_accuracy: 0.4805
Epoch 37/100
504/504 [=====] - ETA: 0s - loss: 0.7083 - accuracy: 0.7360
Epoch 37: val_loss did not improve from 0.97295
504/504 [=====] - 29s 57ms/step - loss: 0.7083 - accuracy: 0.7360 - val_loss: 1.0107 - val_accuracy: 0.6343
Epoch 38/100
504/504 [=====] - ETA: 0s - loss: 0.6990 - accuracy: 0.7375
Epoch 38: val_loss did not improve from 0.97295
504/504 [=====] - 29s 58ms/step - loss: 0.6990 - accuracy: 0.7375 - val_loss: 0.9897 - val_accuracy: 0.6399
Epoch 39/100
504/504 [=====] - ETA: 0s - loss: 0.6908 - accuracy: 0.7412
Epoch 39: val_loss did not improve from 0.97295
504/504 [=====] - 29s 57ms/step - loss: 0.6908 - accuracy: 0.7412 - val_loss: 1.0663 - val_accuracy: 0.6332
Epoch 40/100
504/504 [=====] - ETA: 0s - loss: 0.6807 - accuracy: 0.7446
Epoch 40: val_loss did not improve from 0.97295
504/504 [=====] - 29s 58ms/step - loss: 0.6807 - accuracy: 0.7446 - val_loss: 1.0874 - val_accuracy: 0.6221
Epoch 41/100
504/504 [=====] - ETA: 0s - loss: 0.6735 - accuracy: 0.7472
Epoch 41: val_loss did not improve from 0.97295

```


```
1 score = facialmodel.evaluate(X_test, y_test)
2 print('Test Accuracy: {}'.format(score[1]))
```

57/57 [=====] - 0s 5ms/step - loss: 0.9472 - accuracy: 0.6819
Test Accuracy: 0.6818941235542297

```
1 accuracy = history.history['accuracy']
2 val_accuracy = history.history['val_accuracy']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']
```


```
1 epochs = range(len(accuracy))
2
3 plt.plot(epochs, accuracy, 'bo', label='Training Accuracy')
4 plt.plot(epochs, val_accuracy, 'b', label='Validation Accuracy')
5 plt.title('Training and Validation Accuracy')
6 plt.legend()
```

<matplotlib.legend.Legend at 0x7b0eaad13040>



```
1 plt.plot(epochs, loss, 'ro', label='Training loss')
2 plt.plot(epochs, val_loss, 'r', label='Validation loss')
3 plt.title('Training and Validation loss')
4 plt.legend()
```

<matplotlib.legend.Legend at 0x7b0e241be620>




```
1 predicted_classes = np.argmax(facialmodel.predict(X_test), axis=-1)
2 y_true = np.argmax(y_test, axis=-1)
```

57/57 [=====] - 1s 3ms/step

```
1 from sklearn.metrics import confusion_matrix
2 cm = confusion_matrix(y_true, predicted_classes)
3 plt.figure(figsize = (10, 10))
4 sns.heatmap(cm, annot = True, cbar = False)
```

<Axes: >



```
1 L = 5
2 W = 5
3
4 fig, axes = plt.subplots(L, W, figsize = (24, 24))
5 axes = axes.ravel()
6
7 for i in np.arange(0, L*W):
8     axes[i].imshow(X_test[i].reshape(96,96), cmap = 'gray')
9     axes[i].set_title('Prediction = {}\n True = {}'.format(label_to_text[predicted_classes[i]], label_to_text[y_true[i]]))
10    axes[i].axis('off')
11
12 plt.subplots_adjust(wspace = 1)
```




Prediction = happiness
True = happiness



Prediction = anger
True = anger



Prediction = happiness
True = fear



Prediction = neutral
True = neutral



Prediction = happiness
True = happiness



Prediction = neutral
True = fear



Prediction = sad
True = neutral



Prediction = sad
True = sad



Prediction = sad
True = sad



Prediction = neutral
True = neutral



Prediction = anger
True = anger




Prediction = neutral
True = neutral



Prediction = fear
True = sad



Prediction = surprise
True = surprise



Prediction = fear
True = fear



Prediction = fear
True = fear



Prediction = surprise
True = surprise



Prediction = happiness
True = happiness



Prediction = surprise
True = neutral



Prediction = surprise
True = surprise



Prediction = anger
True = anger



Prediction = happiness
True = fear



Prediction = happiness
True = happiness



Prediction = happiness
True = happiness




Prediction = happiness
True = happiness



1 Start coding or [generate](#) with AI.

```
1 # Compute accuracy
2 accuracy = np.mean(predicted_classes == y_true)
3
4 print("Accuracy:", accuracy)
```

 Accuracy: 0.6818941504178273

```

1 def identity_block(X, f, filters, stage, block):
2     """
3     Implementation of the identity block as defined in Figure 3
4
5     Arguments:
6     X -- input tensor of shape (m, n_H_prev, n_W_prev, n_C_prev)
7     f -- integer, specifying the shape of the middle CONV's window for the main path
8     filters -- python list of integers, defining the number of filters in the CONV layers of the main path
9     stage -- integer, used to name the layers, depending on their position in the network
10    block -- string/character, used to name the layers, depending on their position in the network
11
12    Returns:
13    X -- output of the identity block, tensor of shape (n_H, n_W, n_C)
14    """
15
16    # defining name basis
17    conv_name_base = 'res' + str(stage) + block + '_branch'
18    bn_name_base = 'bn' + str(stage) + block + '_branch'
19
20    # Retrieve Filters
21    F1, F2, F3 = filters
22
23    # Save the input value. You'll need this later to add back to the main path.
24    X_shortcut = X
25
26    # First component of main path
27    X = Conv2D(F1, (1, 1), strides=(1, 1), name=conv_name_base + '2a', kernel_initializer=glorot_uniform(seed=0))(X)
28    X = BatchNormalization(axis=3, name=bn_name_base + '2a')(X)
29    X = Activation('relu')(X)
30
31    # Second component of main path (≈3 lines)
32    X = Conv2D(F2, (f, f), strides=(1, 1), padding='same', name=conv_name_base + '2b', kernel_initializer=glorot_uniform(seed=0))(X)
33    X = BatchNormalization(axis=3, name=bn_name_base + '2b')(X)
34    X = Activation('relu')(X)
35
36    # Third component of main path (≈2 lines)
37    X = Conv2D(F3, (1, 1), strides=(1, 1), name=conv_name_base + '2c', kernel_initializer=glorot_uniform(seed=0))(X)
38    X = BatchNormalization(axis=3, name=bn_name_base + '2c')(X)
39
40    # Add shortcut value to main path, and pass it through a RELU activation (≈2 lines)
41    X = Add()([X, X_shortcut])
42    X = Activation('relu')(X)
43
44    return X
45
46 def convolutional_block(X, f, filters, stage, block, s=2):
47     """
48     Implementation of the convolutional block as defined in Figure 4
49
50     Arguments:
51     X -- input tensor of shape (m, n_H_prev, n_W_prev, n_C_prev)
52     f -- integer, specifying the shape of the middle CONV's window for the main path
53     filters -- python list of integers, defining the number of filters in the CONV layers of the main path
54     stage -- integer, used to name the layers, depending on their position in the network
55     block -- string/character, used to name the layers, depending on their position in the network
56     s -- Integer, specifying the stride to be used
57
58     Returns:
59     X -- output of the convolutional block, tensor of shape (n_H, n_W, n_C)
60     """
61
62    # defining name basis
63    conv_name_base = 'res' + str(stage) + block + '_branch'
64    bn_name_base = 'bn' + str(stage) + block + '_branch'
65
66    # Retrieve Filters
67    F1, F2, F3 = filters
68
69    # Save the input value
70    X_shortcut = X
71
72    # First component of main path
73    X = Conv2D(F1, (1, 1), strides=(s, s), name=conv_name_base + '2a', kernel_initializer=glorot_uniform(seed=0))(X)
74    X = BatchNormalization(axis=3, name=bn_name_base + '2a')(X)
75    X = Activation('relu')(X)
76
77    # Second component of main path

```



```

78 X = Conv2D(F2, (f, f), strides=(1, 1), padding='same', name=conv_name_base + '2b', kernel_initializer=glorot_uniform(seed=0))(X)
79 X = BatchNormalization(axis=3, name=bn_name_base + '2b')(X)
80 X = Activation('relu')(X)
81
82 # Third component of main path
83 X = Conv2D(F3, (1, 1), strides=(1, 1), name=conv_name_base + '2c', kernel_initializer=glorot_uniform(seed=0))(X)
84 X = BatchNormalization(axis=3, name=bn_name_base + '2c')(X)
85
86 # Shortcut path
87 X_shortcut = Conv2D(F3, (1, 1), strides=(s, s), name=conv_name_base + '1', kernel_initializer=glorot_uniform(seed=0))(X_shortcut)
88 X_shortcut = BatchNormalization(axis=3, name=bn_name_base + '1')(X_shortcut)
89
90 # Add shortcut value to main path
91 X = Add()([X, X_shortcut])
92 X = Activation('relu')(X)
93
94 return X
95
96 def ResNet50(input_shape=(96, 96, 1), classes=7):
97     """
98     Implementation of the popular ResNet50 the following architecture:
99     CONV2D -> BATCHNORM -> RELU -> MAXPOOL -> CONVBLOCK -> IDBLOCK*2 -> CONVBLOCK -> IDBLOCK*3
100     -> CONVBLOCK -> IDBLOCK*5 -> CONVBLOCK -> IDBLOCK*2 -> AVGPPOOL -> TOPPLAYER
101
102     Arguments:
103     input_shape -- shape of the images of the dataset
104     classes -- integer, number of classes
105
106     Returns:
107     model -- a Model() instance in Keras
108     """
109
110     # Define the input as a tensor with shape input_shape
111     X_input = Input(input_shape)
112
113     # Zero-Padding
114     X = ZeroPadding2D((3, 3))(X_input)
115
116     # Stage 1
117     X = Conv2D(64, (7, 7), strides=(2, 2), name='conv1', kernel_initializer=glorot_uniform(seed=0))(X)
118     X = BatchNormalization(axis=3, name='bn_conv1')(X)
119     X = Activation('relu')(X)
120     X = MaxPooling2D((3, 3), strides=(2, 2))(X)
121
122     # Stage 2
123     X = convolutional_block(X, f=3, filters=[64, 64, 256], stage=2, block='a', s=1)
124     X = identity_block(X, 3, [64, 64, 256], stage=2, block='b')
125     X = identity_block(X, 3, [64, 64, 256], stage=2, block='c')
126
127     # Stage 3
128     X = convolutional_block(X, f=3, filters=[128, 128, 512], stage=3, block='a', s=2)
129     X = identity_block(X, 3, [128, 128, 512], stage=3, block='b')
130     X = identity_block(X, 3, [128, 128, 512], stage=3, block='c')
131     X = identity_block(X, 3, [128, 128, 512], stage=3, block='d')
132
133     # Stage 4
134     X = convolutional_block(X, f=3, filters=[256, 256, 1024], stage=4, block='a', s=2)
135     X = identity_block(X, 3, [256, 256, 1024], stage=4, block='b')
136     X = identity_block(X, 3, [256, 256, 1024], stage=4, block='c')
137     X = identity_block(X, 3, [256, 256, 1024], stage=4, block='d')
138     X = identity_block(X, 3, [256, 256, 1024], stage=4, block='e')
139     X = identity_block(X, 3, [256, 256, 1024], stage=4, block='f')
140
141     # Stage 5
142     X = convolutional_block(X, f=3, filters=[512, 512, 2048], stage=5, block='a', s=2)
143     X = identity_block(X, 3, [512, 512, 2048], stage=5, block='b')
144     X = identity_block(X, 3, [512, 512, 2048], stage=5, block='c')
145
146     # AVGPPOOL (≈1 line). Use "X = AveragePooling2D(...)(X)"
147     X = AveragePooling2D(pool_size=(2, 2), padding='same')(X)
148
149     # output layer
150     X = Flatten()(X)
151     X = Dense(classes, activation='softmax', name='fc' + str(classes), kernel_initializer=glorot_uniform(seed=0))(X)
152
153     # Create model
154     model = Model(inputs=X_input, outputs=X, name='ResNet50')

```

```

155
156     return model
157
158 # Create ResNet50 model
159 facial_model = ResNet50(input_shape=(96, 96, 1), classes=7)
160
161 # Print model summary
162 facial_model.summary()

```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 96, 96, 1)]	0
block1_conv1 (Conv2D)	(None, 96, 96, 64)	640
block1_conv2 (Conv2D)	(None, 96, 96, 64)	36928
block1_pool (MaxPooling2D)	(None, 48, 48, 64)	0
block2_conv1 (Conv2D)	(None, 48, 48, 128)	73856
block2_conv2 (Conv2D)	(None, 48, 48, 128)	147584
block2_pool (MaxPooling2D)	(None, 24, 24, 128)	0
block3_conv1 (Conv2D)	(None, 24, 24, 256)	295168
block3_conv2 (Conv2D)	(None, 24, 24, 256)	590080
block3_conv3 (Conv2D)	(None, 24, 24, 256)	590080
block3_pool (MaxPooling2D)	(None, 12, 12, 256)	0
block4_conv1 (Conv2D)	(None, 12, 12, 512)	1180160
block4_conv2 (Conv2D)	(None, 12, 12, 512)	2359808
block4_conv3 (Conv2D)	(None, 12, 12, 512)	2359808
block4_pool (MaxPooling2D)	(None, 6, 6, 512)	0
block5_conv1 (Conv2D)	(None, 6, 6, 512)	2359808
block5_conv2 (Conv2D)	(None, 6, 6, 512)	2359808
block5_conv3 (Conv2D)	(None, 6, 6, 512)	2359808
block5_pool (MaxPooling2D)	(None, 3, 3, 512)	0
flatten (Flatten)	(None, 4608)	0
fc1 (Dense)	(None, 4096)	18878464
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 7)	28679
Total params: 50401991 (192.27 MB)		
Trainable params: 50401991 (192.27 MB)		
Non-trainable params: 0 (0.00 Byte)		

```

1
2
3 # Compile the model
4 facial_model.compile(optimizer=keras.optimizers.Adam(learning_rate=1e-3), loss = "categorical_crossentropy", metrics = ["accuracy"])

1 # using early stopping to exit training if validation loss is not decreasing even after certain epochs (patience)
2 earlystopping = EarlyStopping(monitor = 'val_loss', mode = 'min', verbose = 1, patience = 20)
3
4 # save the best model with lower validation loss
5 checkpointer = ModelCheckpoint(filepath = "/content/drive/MyDrive/DL Facial Recognition /FacialExpression_weights_VGG.hdf5", verbose = 1,

1 history = facial_model.fit(train_datagen.flow(X_train, y_train, batch_size=64),
2     validation_data=(X_val, y_val), steps_per_epoch=len(X_train) // 64,
3     epochs= 100, callbacks=[checkpointer, earlystopping])

```

```

Epoch 1/100
504/504 [=====] - ETA: 0s - loss: 1.8179 - accuracy: 0.2489
Epoch 1: val_loss improved from inf to 1.80318, saving model to /content/drive/MyDrive/DL Facial Recognition /FacialExpression_weights
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via
saving_api.save_model(
504/504 [=====] - 40s 65ms/step - loss: 1.8179 - accuracy: 0.2489 - val_loss: 1.8032 - val_accuracy: 0.2469
Epoch 2/100
504/504 [=====] - ETA: 0s - loss: 1.8131 - accuracy: 0.2504
Epoch 2: val_loss improved from 1.80318 to 1.80172, saving model to /content/drive/MyDrive/DL Facial Recognition /FacialExpression_weights
504/504 [=====] - 36s 71ms/step - loss: 1.8131 - accuracy: 0.2504 - val_loss: 1.8017 - val_accuracy: 0.2469
Epoch 3/100
504/504 [=====] - ETA: 0s - loss: 1.8127 - accuracy: 0.2502
Epoch 3: val_loss improved from 1.80172 to 1.80084, saving model to /content/drive/MyDrive/DL Facial Recognition /FacialExpression_weights
504/504 [=====] - 36s 70ms/step - loss: 1.8127 - accuracy: 0.2502 - val_loss: 1.8008 - val_accuracy: 0.2469
Epoch 4/100
504/504 [=====] - ETA: 0s - loss: 1.8128 - accuracy: 0.2504
Epoch 4: val_loss did not improve from 1.80084
504/504 [=====] - 29s 58ms/step - loss: 1.8128 - accuracy: 0.2504 - val_loss: 1.8027 - val_accuracy: 0.2469
Epoch 5/100
504/504 [=====] - ETA: 0s - loss: 1.8125 - accuracy: 0.2504
Epoch 5: val_loss did not improve from 1.80084
504/504 [=====] - 29s 57ms/step - loss: 1.8125 - accuracy: 0.2504 - val_loss: 1.8039 - val_accuracy: 0.2469
Epoch 6/100
504/504 [=====] - ETA: 0s - loss: 1.8125 - accuracy: 0.2504
Epoch 6: val_loss did not improve from 1.80084
504/504 [=====] - 29s 58ms/step - loss: 1.8125 - accuracy: 0.2504 - val_loss: 1.8016 - val_accuracy: 0.2469
Epoch 7/100
504/504 [=====] - ETA: 0s - loss: 1.8126 - accuracy: 0.2504
Epoch 7: val_loss did not improve from 1.80084
504/504 [=====] - 29s 58ms/step - loss: 1.8126 - accuracy: 0.2504 - val_loss: 1.8017 - val_accuracy: 0.2469
Epoch 8/100
504/504 [=====] - ETA: 0s - loss: 1.8123 - accuracy: 0.2505
Epoch 8: val_loss did not improve from 1.80084
504/504 [=====] - 29s 57ms/step - loss: 1.8123 - accuracy: 0.2505 - val_loss: 1.8022 - val_accuracy: 0.2469
Epoch 9/100
504/504 [=====] - ETA: 0s - loss: 1.8121 - accuracy: 0.2504
Epoch 9: val_loss did not improve from 1.80084
504/504 [=====] - 29s 58ms/step - loss: 1.8121 - accuracy: 0.2504 - val_loss: 1.8013 - val_accuracy: 0.2469
Epoch 10/100
504/504 [=====] - ETA: 0s - loss: 1.8120 - accuracy: 0.2504
Epoch 10: val_loss did not improve from 1.80084
504/504 [=====] - 29s 58ms/step - loss: 1.8120 - accuracy: 0.2504 - val_loss: 1.8012 - val_accuracy: 0.2469
Epoch 11/100
504/504 [=====] - ETA: 0s - loss: 1.8120 - accuracy: 0.2504
Epoch 11: val_loss improved from 1.80084 to 1.80050, saving model to /content/drive/MyDrive/DL Facial Recognition /FacialExpression_weights
504/504 [=====] - 36s 71ms/step - loss: 1.8120 - accuracy: 0.2504 - val_loss: 1.8005 - val_accuracy: 0.2469
Epoch 12/100
504/504 [=====] - ETA: 0s - loss: 1.8118 - accuracy: 0.2506
Epoch 12: val_loss did not improve from 1.80050
504/504 [=====] - 29s 58ms/step - loss: 1.8118 - accuracy: 0.2506 - val_loss: 1.8008 - val_accuracy: 0.2469
Epoch 13/100
504/504 [=====] - ETA: 0s - loss: 1.8118 - accuracy: 0.2504
Epoch 13: val_loss did not improve from 1.80050
504/504 [=====] - 29s 57ms/step - loss: 1.8118 - accuracy: 0.2504 - val_loss: 1.8018 - val_accuracy: 0.2469
Epoch 14/100
504/504 [=====] - ETA: 0s - loss: 1.8119 - accuracy: 0.2503
Epoch 14: val_loss did not improve from 1.80050

```

```

1 score = facial_model.evaluate(X_test, y_test)
2 print('Test Accuracy: {}'.format(score[1]))

```

```

57/57 [=====] - 1s 12ms/step - loss: 1.8113 - accuracy: 0.2552
Test Accuracy: 0.25515320897102356


```

```

1 accuracy = history.history['accuracy']
2 val_accuracy = history.history['val_accuracy']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']

1 epochs = range(len(accuracy))
2
3 plt.plot(epochs, accuracy, 'bo', label='Training Accuracy')
4 plt.plot(epochs, val_accuracy, 'b', label='Validation Accuracy')
5 plt.title('Training and Validation Accuracy')
6 plt.legend()

```

 <matplotlib.legend.Legend at 0x7b0d40dc3700>

