

```

1 import pandas as pd
2 import seaborn as sns
3 import numpy as np
4 import os
5 import PIL
6 import pickle
7 from PIL import *
8 import cv2
9 import tensorflow as tf
10 import keras
11 from keras.applications import DenseNet121 # 2017 architecture
12 from keras.models import Model, load_model
13 from keras.initializers import glorot_uniform
14 from keras.utils import plot_model
15 from keras.callbacks import ReduceLROnPlateau, EarlyStopping, ModelCheckpoint, LearningRateScheduler
16 from IPython.display import display
17 from keras import *
18 from keras.preprocessing.image import ImageDataGenerator
19 from keras import layers, optimizers
20 from keras.applications.resnet50 import ResNet50
21 from keras.layers import *
22 from keras import backend as K
23 from keras import optimizers
24 import matplotlib.pyplot as plt
25 from sklearn.model_selection import train_test_split

```

```

1 from google.colab import drive
2 drive.mount('/content/drive')
3 facialexpression_df = pd.read_csv('/content/drive/MyDrive/DL Facial Recognition/icml_face_data.csv')
4

```

↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```

1 facialexpression_df = facialexpression_df.drop(columns = " Usage")
2

```

```
1 facialexpression_df.head()
```

↗

	emotion	pixels
0	0	70 80 82 72 58 58 60 63 54 58 60 48 89 115 121...
1	0	151 150 147 155 148 133 111 140 170 174 182 15...
2	2	231 212 156 164 174 138 161 173 182 200 106 38...
3	4	24 32 36 30 32 23 19 20 30 41 21 22 32 34 21 1...
4	6	4 0 0 0 0 0 0 0 0 0 0 0 3 15 23 28 48 50 58 84...

```

1 def string2array(x):
2     return np.array(x.split(' ')).reshape(48,48,1).astype('float32')

```

```

1 def resize(x):
2     img = x.reshape(48,48)
3     temp = cv2.resize(img, dsize=(48,48), interpolation = cv2.INTER_CUBIC)
4     return np.stack([temp,temp,temp], axis = 2)

```

```

1 facialexpression_df[' pixels'] = facialexpression_df[' pixels'].apply(lambda x : string2array(x))
2 facialexpression_df[' pixels'] = facialexpression_df[' pixels'].apply(lambda x : resize(x))

```

```
1 facialexpression_df.shape
```

↗ (35887, 2)

```
1 label_to_text = {0: 'anger', 1 : 'disgust', 2 : 'fear', 3 : 'happiness', 4: 'sad', 5: 'surprise', 6: 'neutral'}
```

```
1 facialexpression_df[' pixels'][1].shape
```

↗ (48, 48, 3)

```

1 emotions = [1]
2 for i in emotions:
3     data = facialexpression_df[facialexpression_df['emotion']==i][:1]
4
5     img = data[' pixels'].item()
6     img = img.reshape(48,48, 3)
7
8     plt.figure()
9     plt.title(label_to_text[i])
10    plt.imshow(img/255)
11    plt.axis('off')

```



disgust



```
1 facialexpression_df.emotion.value_counts().index
```



```
Index([3, 6, 4, 2, 0, 5, 1], dtype='int64', name='emotion')
```

```
1 facialexpression_df.emotion.value_counts()
```



```

emotion
3      8989
6      6198
4      6077
2      5121
0      4953
5      4002
1       547
Name: count, dtype: int64

```

```

1 class1 = facialexpression_df.loc[facialexpression_df['emotion'] == 1]
2 class1.shape

```



```
(547, 2)
```

```

1 from keras.utils import to_categorical
2
3 X = facialexpression_df[' pixels']
4 y = to_categorical(facialexpression_df['emotion'])
5
6 X = np.stack(X, axis = 0)
7 X = X.reshape(35887,48,48,3)
8
9 print(X.shape, y.shape)

```



```
(35887, 48, 48, 3) (35887, 7)
```

```

1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, shuffle = True)
3 X_val, X_test, y_val, y_test = train_test_split(X_test, y_test, test_size = 0.5, shuffle = True)

```

```
1 X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

→ ((32298, 48, 48, 3), (1795, 48, 48, 3), (32298, 7), (1795, 7))

```
1 X_train = X_train/255
2 X_val   = X_val /255
3 X_test  = X_test/255
```

```
1 train_datagen = ImageDataGenerator(
2     rotation_range = 15,
3     width_shift_range = 0.1,
4     height_shift_range = 0.1,
5     shear_range = 0.1,
6     zoom_range = 0.1,
7     horizontal_flip = True,
8     fill_mode = "nearest")
```

```
1 from keras.applications.resnet50 import ResNet50
2
3 input_shape = (48,48, 3)
4
5 base_model = ResNet50(weights='imagenet', include_top=False, input_shape=input_shape)
6 for layer in base_model.layers:
7     layer.trainable = True
8     if isinstance(layer, tf.keras.layers.BatchNormalization):
9         layer.trainable = True
10 Resnet = Sequential([
11     base_model,
12     Flatten(),
13     Dense(256, activation='relu'),
14     Dense(7, activation='sigmoid')
15 ])
16
17 Resnet.summary()
```

→ Model: "sequential_3"

Layer (type)	Output Shape	Param #
resnet50 (Functional)	(None, 2, 2, 2048)	23587712
flatten_3 (Flatten)	(None, 8192)	0
dense_6 (Dense)	(None, 256)	2097408
dense_7 (Dense)	(None, 7)	1799
Total params: 25686919 (97.99 MB)		
Trainable params: 25633799 (97.79 MB)		
Non-trainable params: 53120 (207.50 KB)		

```
1 reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=10, verbose=1, min_lr=1e-6)
2 Resnet.compile(optimizer=keras.optimizers.Adam(learning_rate=1e-3), loss = "categorical_crossentropy", metrics = ["accuracy"])

1 # using early stopping to exit training if validation loss is not decreasing even after certain epochs (patience)
2 earlystopping = EarlyStopping(monitor = 'val_loss', mode = 'min', verbose = 1, patience = 20)
3
4 # save the best model with lower validation loss
5 checkpointer = ModelCheckpoint(filepath = "/content/drive/MyDrive/DL Facial Recognition /Resnet50_weights.hdf5", verbose = 1, save_best_c

1 history = Resnet.fit(train_datagen.flow(X_train, y_train, batch_size=64),
2     validation_data= (X_val, y_val), steps_per_epoch=len(X_train) // 64,
3     epochs= 100, callbacks=[checkpointer, earlystopping, reduce_lr])
```

→

```

Epoch 72/100
504/504 [=====] - ETA: 0s - loss: 1.0303 - accuracy: 0.6036
Epoch 72: val_loss did not improve from 1.03164
504/504 [=====] - 29s 58ms/step - loss: 1.0303 - accuracy: 0.6036 - val_loss: 1.0562 - val_accuracy: 0.5897
Epoch 73/100
504/504 [=====] - ETA: 0s - loss: 1.0256 - accuracy: 0.6071
Epoch 73: val_loss did not improve from 1.03164
504/504 [=====] - 29s 58ms/step - loss: 1.0256 - accuracy: 0.6071 - val_loss: 1.0761 - val_accuracy: 0.5842
Epoch 74/100
504/504 [=====] - ETA: 0s - loss: 1.0200 - accuracy: 0.6078
Epoch 74: val_loss did not improve from 1.03164

Epoch 74: ReduceLROnPlateau reducing learning rate to 1.0000000474974514e-05.
504/504 [=====] - 29s 58ms/step - loss: 1.0200 - accuracy: 0.6078 - val_loss: 1.0550 - val_accuracy: 0.5892
Epoch 75/100
504/504 [=====] - ETA: 0s - loss: 1.0160 - accuracy: 0.6100
Epoch 75: val_loss did not improve from 1.03164
504/504 [=====] - 29s 58ms/step - loss: 1.0160 - accuracy: 0.6100 - val_loss: 1.0575 - val_accuracy: 0.5864
Epoch 76/100
504/504 [=====] - ETA: 0s - loss: 1.0166 - accuracy: 0.6103
Epoch 76: val_loss did not improve from 1.03164
504/504 [=====] - 29s 58ms/step - loss: 1.0166 - accuracy: 0.6103 - val_loss: 1.0573 - val_accuracy: 0.5808
Epoch 77/100
504/504 [=====] - ETA: 0s - loss: 1.0192 - accuracy: 0.6072
Epoch 77: val_loss did not improve from 1.03164
504/504 [=====] - 29s 58ms/step - loss: 1.0192 - accuracy: 0.6072 - val_loss: 1.0547 - val_accuracy: 0.5836
Epoch 78/100
504/504 [=====] - ETA: 0s - loss: 1.0155 - accuracy: 0.6109
Epoch 78: val_loss did not improve from 1.03164
504/504 [=====] - 29s 58ms/step - loss: 1.0155 - accuracy: 0.6109 - val_loss: 1.0556 - val_accuracy: 0.5836
Epoch 79/100
504/504 [=====] - ETA: 0s - loss: 1.0172 - accuracy: 0.6086
Epoch 79: val_loss did not improve from 1.03164
504/504 [=====] - 30s 59ms/step - loss: 1.0172 - accuracy: 0.6086 - val_loss: 1.0563 - val_accuracy: 0.5847
Epoch 80/100
504/504 [=====] - ETA: 0s - loss: 1.0137 - accuracy: 0.6099
Epoch 80: val_loss did not improve from 1.03164
504/504 [=====] - 29s 58ms/step - loss: 1.0137 - accuracy: 0.6099 - val_loss: 1.0577 - val_accuracy: 0.5825
Epoch 81/100
504/504 [=====] - ETA: 0s - loss: 1.0186 - accuracy: 0.6065
Epoch 81: val_loss did not improve from 1.03164
504/504 [=====] - 29s 58ms/step - loss: 1.0186 - accuracy: 0.6065 - val_loss: 1.0560 - val_accuracy: 0.5831
Epoch 82/100
504/504 [=====] - ETA: 0s - loss: 1.0149 - accuracy: 0.6115
Epoch 82: val_loss did not improve from 1.03164
504/504 [=====] - 29s 58ms/step - loss: 1.0149 - accuracy: 0.6115 - val_loss: 1.0563 - val_accuracy: 0.5810

```

```

1 score = Resnet.evaluate(X_test, y_test)
2 print('Test Accuracy: {}'.format(score[1]))

```

```

57/57 [=====] - 0s 8ms/step - loss: 1.0112 - accuracy: 0.5961
Test Accuracy: 0.5961002707481384

```

```

1 accuracy = history.history['accuracy']
2 val_accuracy = history.history['val_accuracy']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']


```

```

1 epochs = range(len(accuracy))
2
3 plt.plot(epochs, accuracy, 'bo', label='Training Accuracy')
4 plt.plot(epochs, val_accuracy, 'b', label='Validation Accuracy')
5 plt.title('Training and Validation Accuracy')
6 plt.legend()


```

 <matplotlib.legend.Legend at 0x7c1d4854efb0>




```
1 plt.plot(epochs, loss, 'ro', label='Training loss')
2 plt.plot(epochs, val_loss, 'r', label='Validation loss')
3 plt.title('Training and Validation loss')
4 plt.legend()
```



 <matplotlib.legend.Legend at 0x7c1d4847bdf0>



```
1 predicted_classes = np.argmax(Resnet.predict(X_test), axis=-1)
2 y_true = np.argmax(y_test, axis=-1)
```

 57/57 [=====] - 0s 7ms/step

```
1 from sklearn.metrics import confusion_matrix
2 cm = confusion_matrix(y_true, predicted_classes)
3 plt.figure(figsize = (10, 10))
4 sns.heatmap(cm, annot = True, cbar = False)
```

 <Axes: >


```

1 L = 5
2 W = 5
3
4 fig, axes = plt.subplots(L, W, figsize = (24, 24))
5 axes = axes.ravel()
6
7 for i in np.arange(0, L*W):
8     axes[i].imshow(X_test[i], cmap = 'gray')
9     axes[i].set_title('Prediction = {}\n True = {}'.format(label_to_text[predicted_classes[i]], label_to_text[y_true[i]]))
10    axes[i].axis('off')
11
12 plt.subplots_adjust(wspace = 1)

```



<div>Prediction = sad True = fear</div>	<div>Prediction = neutral True = fear</div>	<div>Prediction = neutral True = anger</div>	<div>Prediction = happiness True = happiness</div>	<div>Prediction = happiness True = sad</div>
<div>Prediction = happiness True = happiness</div>	<div>Prediction = happiness True = neutral</div>	<div>Prediction = neutral True = anger</div>	<div>Prediction = surprise True = surprise</div>	<div>Prediction = surprise True = surprise</div>
<div>Prediction = anger True = happiness</div>	<div>Prediction = neutral True = neutral</div>	<div>Prediction = surprise True = surprise</div>	<div>Prediction = neutral True = neutral</div>	<div>Prediction = neutral True = fear</div>
<div>Prediction = neutral True = neutral</div>	<div>Prediction = sad True = neutral</div>	<div>Prediction = happiness True = happiness</div>	<div>Prediction = anger True = happiness</div>	<div>Prediction = happiness True = happiness</div>
<div>Prediction = neutral True = sad</div>	<div>Prediction = neutral True = fear</div>	<div>Prediction = neutral True = sad</div>	<div>Prediction = happiness True = happiness</div>	<div>Prediction = anger True = sad</div>

```
1 # Compute accuracy
```