


```
1 import pandas as pd
2 import seaborn as sns
3 import numpy as np
4 import os
5 import PIL
6 import pickle
7 from PIL import *
8 import cv2
9 import tensorflow as tf
10 import keras
11 from keras.applications import DenseNet121 # 2017 architecture
12 from keras.models import Model, load_model
13 from keras.initializers import glorot_uniform
14 from keras.utils import plot_model
15 from keras.callbacks import ReduceLROnPlateau, EarlyStopping, ModelCheckpoint, LearningRateScheduler
16 from IPython.display import display
17 from keras import *
18 from keras.preprocessing.image import ImageDataGenerator
19 from keras import layers, optimizers
20 from keras.applications.resnet50 import ResNet50
21 from keras.layers import *
22 from keras import backend as K
23 from keras import optimizers
24 import matplotlib.pyplot as plt
25 from sklearn.model_selection import train_test_split
26 from keras.regularizers import l2
27
```

```
1 from google.colab import drive
2 drive.mount('/content/drive')
3
4 # Load the dataset
5 facialexpression_df = pd.read_csv('/content/drive/MyDrive/DL Facial Recognition/icml_face_data.csv')
6 facialexpression_df = facialexpression_df.drop(columns=" Usage")
```

 Mounted at /content/drive

```

1 def conv_block(X, f, filters, stage, block, strides=(2, 2), weight_decay=1e-3):
2     F1, F2, F3 = filters
3     conv_name_base = 'res' + str(stage) + block + '_branch'
4     bn_name_base = 'bn' + str(stage) + block + '_branch'
5
6     # Save the input value for shortcut
7     X_shortcut = X
8
9     # First component of main path
10    X = Conv2D(F1, (1, 1), strides=strides, padding='valid', name=conv_name_base + '2a',
11              kernel_initializer=glorot_uniform(seed=0), kernel_regularizer=l2(weight_decay))(X)
12    X = BatchNormalization(axis=3, name=bn_name_base + '2a')(X)
13    X = Activation('relu')(X)
14
15    # Second component
16    X = Conv2D(F2, (f, f), padding='same', name=conv_name_base + '2b',
17              kernel_initializer=glorot_uniform(seed=0), kernel_regularizer=l2(weight_decay))(X)
18    X = BatchNormalization(axis=3, name=bn_name_base + '2b')(X)
19    X = Activation('relu')(X)
20
21    # Third component
22    X = Conv2D(F3, (1, 1), padding='valid', name=conv_name_base + '2c',
23              kernel_initializer=glorot_uniform(seed=0), kernel_regularizer=l2(weight_decay))(X)
24    X = BatchNormalization(axis=3, name=bn_name_base + '2c')(X)
25
26    # Shortcut path
27    X_shortcut = Conv2D(F3, (1, 1), strides=strides, padding='valid', name=conv_name_base + '1',
28                      kernel_initializer=glorot_uniform(seed=0), kernel_regularizer=l2(weight_decay))(X_shortcut)
29    X_shortcut = BatchNormalization(axis=3, name=bn_name_base + '1')(X_shortcut)
30
31    # Final step: Add shortcut value to main path, and pass it through a RELU activation
32    X = Add()([X, X_shortcut])
33    X = Activation('relu')(X)
34
35    return X
36
37 def identity_block(X, f, filters, stage, block):
38     """The identity block is the block that has no conv layer at shortcut."""
39     F1, F2, F3 = filters
40     conv_name_base = 'res' + str(stage) + block + '_branch'
41     bn_name_base = 'bn' + str(stage) + block + '_branch'
42
43     X_shortcut = X
44
45     # First component of main path
46    X = Conv2D(F1, (1, 1), padding='valid', name=conv_name_base + '2a', kernel_initializer=glorot_uniform(seed=0))(X)
47    X = BatchNormalization(axis=3, name=bn_name_base + '2a')(X)
48    X = Activation('relu')(X)
49
50    # Second component
51    X = Conv2D(F2, (f, f), padding='same', name=conv_name_base + '2b', kernel_initializer=glorot_uniform(seed=0))(X)
52    X = BatchNormalization(axis=3, name=bn_name_base + '2b')(X)
53    X = Activation('relu')(X)
54
55    # Third component
56    X = Conv2D(F3, (1, 1), padding='valid', name=conv_name_base + '2c', kernel_initializer=glorot_uniform(seed=0))(X)
57    X = BatchNormalization(axis=3, name=bn_name_base + '2c')(X)
58
59    # Add shortcut
60    X = Add()([X, X_shortcut])
61    X = Activation('relu')(X)
62
63    return X
64
65 def apply_blocks(X, filter_triplets, first_strides, num_identity_blocks, stage):
66     """Apply a convolutional block followed by multiple identity blocks."""
67     # Apply the first convolutional block with stride changes
68    X = conv_block(X, 3, filter_triplets, stage, 'a', strides=first_strides)
69    # Apply subsequent identity blocks
70    for i in range(num_identity_blocks):
71        X = identity_block(X, 3, filter_triplets, stage, chr(98 + i)) # 'b', 'c', etc.
72    return X
73

```

```
1
2 # Define function to convert string to array
3 def string2array(x):
4     return np.array(x.split(' ')).reshape(48, 48, 1).astype('float32')
5 # Define function to resize images
6 def resize(x):
7     img = x.reshape(48, 48)
8     return cv2.resize(img, dsize=(96, 96), interpolation=cv2.INTER_CUBIC)
9

1
2 # Apply preprocessing to pixels column
3 facialexpression_df[' pixels'] = facialexpression_df[' pixels'].apply(lambda x: string2array(x))
4 facialexpression_df[' pixels'] = facialexpression_df[' pixels'].apply(lambda x : resize(x))
5 # Define emotions of interest
6 emotions = [1] # Example: you want to work with emotion label 1
7 label_to_text = {0: 'anger', 1 : 'disgust', 2 : 'fear', 3 : 'happiness', 4: 'sad', 5: 'surprise', 6: 'neutral'}

1 # Plot sample images
2 for i in emotions:
3     data = facialexpression_df[facialexpression_df['emotion'] == i][:1]
4     img = data[' pixels'].item()
5     img = img.reshape(96, 96)
6
7     plt.figure()
8     plt.title(label_to_text[i])
9     plt.imshow(img, cmap='gray')
10    plt.axis('off')
```



disgust



```

1 from keras.utils import to_categorical
2
3 # Prepare input and target data
4 X = facialexpression_df[' pixels']
5 y = to_categorical(facialexpression_df['emotion'])
6
7 X = np.stack(X, axis = 0)
8 X = X.reshape(35887,96,96,1)
9
10 from sklearn.model_selection import train_test_split
11
12 # Split data into training, validation, and test sets
13 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, shuffle=True)
14 X_val, X_test, y_val, y_test = train_test_split(X_test, y_test, test_size=0.5, shuffle=True)
15
16 X_train = X_train/255
17 X_val = X_val /255
18 X_test = X_test/255
19
20 # Data Augmentation
21 train_datagen = ImageDataGenerator(
22     rotation_range = 15,
23     width_shift_range = 0.1,
24     height_shift_range = 0.1,
25     shear_range = 0.1,
26     zoom_range = 0.1,
27     horizontal_flip = True,
28     fill_mode = "nearest")
29
30
31 # Define the input tensor shape
32 input_shape = (96, 96, 1)
33 X_input = Input(input_shape)
34
35 # Initial convolution and max pooling
36 X = ZeroPadding2D((3, 3))(X_input)
37 X = Conv2D(64, (7, 7), strides=(2, 2), name='conv1', kernel_initializer=glorot_uniform(seed=0))(X)
38 X = BatchNormalization(axis=3, name='bn_conv1')(X)
39 X = Dropout(0.5)(X) # Add dropout
40 X = Activation('relu')(X)
41 X = MaxPooling2D((3, 3), strides=(2, 2), padding='same')(X)
42
43 # Helper function to apply a convolutional block and multiple identity blocks
44 def apply_blocks(X, filter_triplets, first_strides, num_identity_blocks, stage):
45     # Apply convolutional block
46     X = conv_block(X, 3, filter_triplets, stage, 'a', strides=first_strides)
47     # Apply identity blocks
48     for i in range(num_identity_blocks):
49         X = identity_block(X, 3, filter_triplets, stage, chr(98 + i)) # 'b', 'c', ..., for identity blocks
50     return X
51
52 # Stage 2
53 X = apply_blocks(X, [64, 64, 256], (1, 1), 2, 2)
54
55 # Stage 3
56 X = apply_blocks(X, [128, 128, 512], (2, 2), 3, 3)
57
58 # Stage 4
59 X = apply_blocks(X, [256, 256, 1024], (2, 2), 5, 4)
60
61 # Stage 5
62 X = apply_blocks(X, [512, 512, 2048], (2, 2), 2, 5)
63
64 # Global Average Pooling and final dense layer
65 X = GlobalAveragePooling2D()(X)
66 X = Dense(7, activation='softmax', name='final_output')(X)
67
68 # Create the model
69 model = Model(inputs=X_input, outputs=X, name='CustomResNet50')
70 # Show the model summary
71 model.summary()

```



res5b_branch2b (Conv2D)	(None, 3, 3, 512)	2359808	['activation_43[0][0]']
bn5b_branch2b (BatchNormalization)	(None, 3, 3, 512)	2048	['res5b_branch2b[0][0]']
activation_44 (Activation)	(None, 3, 3, 512)	0	['bn5b_branch2b[0][0]']
res5b_branch2c (Conv2D)	(None, 3, 3, 2048)	1050624	['activation_44[0][0]']
bn5b_branch2c (BatchNormalization)	(None, 3, 3, 2048)	8192	['res5b_branch2c[0][0]']
add_14 (Add)	(None, 3, 3, 2048)	0	['bn5b_branch2c[0][0]', 'activation_42[0][0]']
activation_45 (Activation)	(None, 3, 3, 2048)	0	['add_14[0][0]']
res5c_branch2a (Conv2D)	(None, 3, 3, 512)	1049088	['activation_45[0][0]']
bn5c_branch2a (BatchNormalization)	(None, 3, 3, 512)	2048	['res5c_branch2a[0][0]']
activation_46 (Activation)	(None, 3, 3, 512)	0	['bn5c_branch2a[0][0]']
res5c_branch2b (Conv2D)	(None, 3, 3, 512)	2359808	['activation_46[0][0]']
bn5c_branch2b (BatchNormalization)	(None, 3, 3, 512)	2048	['res5c_branch2b[0][0]']
activation_47 (Activation)	(None, 3, 3, 512)	0	['bn5c_branch2b[0][0]']
res5c_branch2c (Conv2D)	(None, 3, 3, 2048)	1050624	['activation_47[0][0]']
bn5c_branch2c (BatchNormalization)	(None, 3, 3, 2048)	8192	['res5c_branch2c[0][0]']
add_15 (Add)	(None, 3, 3, 2048)	0	['bn5c_branch2c[0][0]', 'activation_45[0][0]']
activation_48 (Activation)	(None, 3, 3, 2048)	0	['add_15[0][0]']
global_average_pooling2d (GlobalAveragePooling2D)	(None, 2048)	0	['activation_48[0][0]']
final_output (Dense)	(None, 7)	14343	['global_average_pooling2d[0][0]']

```

=====
Total params: 23595783 (90.01 MB)
Trainable params: 23542663 (89.81 MB)
Non-trainable params: 53120 (207.50 KB)

```

```

1 # Define the total number of epochs and the number of warmup epochs
2 # Define the total number of epochs, the number of warmup epochs, and the batch size
3 total_epochs = 100
4 warmup_epochs = 5
5 batch_size = 64
6
7 # Define the initial learning rate and minimum learning rate
8 initial_learning_rate = 1e-3
9 min_learning_rate = 1e-6
10
11 # Calculate the total number of steps and warmup steps
12 total_steps = total_epochs * (len(X_train) // batch_size)
13 warmup_steps = warmup_epochs * (len(X_train) // batch_size)
14 reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=10, verbose=1, min_lr=1e-6)
15
16 # Define the learning rate scheduler with warmup
17 # Create the learning rate scheduler with warmup
18 # def lr_schedule(epoch, lr):
19 #     if epoch < warmup_epochs:
20 #         # Linear warmup
21 #         return initial_learning_rate * (epoch + 1) / warmup_epochs
22 #     else:
23 #         # Cosine annealing with restarts
24 #         cosine_decay = 0.5 * (1 + tf.cos(tf.constant(np.pi) * (epoch - warmup_epochs) / (total_epochs - warmup_epochs)))
25 #         decayed = (initial_learning_rate - min_learning_rate) * cosine_decay + min_learning_rate
26 #         return decayed
27 # using early stopping to exit training if validation loss is not decreasing even after certain epochs (patience)
28 earlystopping = EarlyStopping(monitor = 'val_loss', mode = 'min', verbose = 1, patience = 20)
29
30 # save the best model with lower validation loss
31 checkpointer = ModelCheckpoint(filepath = "/content/drive/MyDrive/DL Facial Recognition /FacialExpression_weights_resnet.hdf5", verbose =
32 # Compile the model
33 model.compile(keras.optimizers.Adam(lr = 1e-3), loss='categorical_crossentropy', metrics=['accuracy'])
34 model.summary()
35 train_generator = train_datagen.flow(X_train, y_train, batch_size=batch_size)
36
37

```

WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g.,tf.keras.optimizers.]
Model: "CustomResNet50"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 96, 96, 1)]	0	[]
zero_padding2d (ZeroPaddin g2D)	(None, 102, 102, 1)	0	['input_1[0][0]']
conv1 (Conv2D)	(None, 48, 48, 64)	3200	['zero_padding2d[0][0]']
bn_conv1 (BatchNormalizati on)	(None, 48, 48, 64)	256	['conv1[0][0]']
dropout (Dropout)	(None, 48, 48, 64)	0	['bn_conv1[0][0]']
activation (Activation)	(None, 48, 48, 64)	0	['dropout[0][0]']
max_pooling2d (MaxPooling2 D)	(None, 24, 24, 64)	0	['activation[0][0]']
res2a_branch2a (Conv2D)	(None, 24, 24, 64)	4160	['max_pooling2d[0][0]']
bn2a_branch2a (BatchNormal ization)	(None, 24, 24, 64)	256	['res2a_branch2a[0][0]']
activation_1 (Activation)	(None, 24, 24, 64)	0	['bn2a_branch2a[0][0]']
res2a_branch2b (Conv2D)	(None, 24, 24, 64)	36928	['activation_1[0][0]']
bn2a_branch2b (BatchNormal ization)	(None, 24, 24, 64)	256	['res2a_branch2b[0][0]']
activation_2 (Activation)	(None, 24, 24, 64)	0	['bn2a_branch2b[0][0]']
res2a_branch2c (Conv2D)	(None, 24, 24, 256)	16640	['activation_2[0][0]']
res2a_branch1 (Conv2D)	(None, 24, 24, 256)	16640	['max_pooling2d[0][0]']
bn2a_branch2c (BatchNormal ization)	(None, 24, 24, 256)	1024	['res2a_branch2c[0][0]']

bn2a_branch1 (BatchNormalization)	(None, 24, 24, 256)	1024	['res2a_branch1[0][0]']
add (Add)	(None, 24, 24, 256)	0	['bn2a_branch2c[0][0]', 'bn2a_branch1[0][0]']
activation_3 (Activation)	(None, 24, 24, 256)	0	['add[0][0]']
res2b_branch2a (Conv2D)	(None, 24, 24, 64)	16448	['activation_3[0][0]']
bn2b_branch2a (BatchNormalization)	(None, 24, 24, 64)	256	['res2b_branch2a[0][0]']
activation_4 (Activation)	(None, 24, 24, 64)	0	['bn2b_branch2a[0][0]']

```
1 history = model.fit(train_generator, steps_per_epoch=len(X_train)//batch_size, epochs=total_epochs, validation_data=(X_val, y_val), valid
2
```

```
504/504 [=====] - ETA: 0s - loss: 1.0369 - accuracy: 0.6308
Epoch 39: val_loss improved from 1.12797 to 1.10169, saving model to /content/drive/MyDrive/DL Facial Recognition /FacialExpression_w
504/504 [=====] - 42s 84ms/step - loss: 1.0369 - accuracy: 0.6308 - val_loss: 1.1017 - val_accuracy: 0.6003 -
Epoch 40/100
504/504 [=====] - ETA: 0s - loss: 1.0262 - accuracy: 0.6326
Epoch 40: val_loss did not improve from 1.10169
504/504 [=====] - 38s 74ms/step - loss: 1.0262 - accuracy: 0.6326 - val_loss: 1.1685 - val_accuracy: 0.5808 -
Epoch 41/100
504/504 [=====] - ETA: 0s - loss: 1.0206 - accuracy: 0.6362
Epoch 41: val_loss did not improve from 1.10169
504/504 [=====] - 37s 74ms/step - loss: 1.0206 - accuracy: 0.6362 - val_loss: 1.1742 - val_accuracy: 0.5747 -
Epoch 42/100
504/504 [=====] - ETA: 0s - loss: 1.0144 - accuracy: 0.6385
Epoch 42: val_loss did not improve from 1.10169
504/504 [=====] - 37s 74ms/step - loss: 1.0144 - accuracy: 0.6385 - val_loss: 1.1298 - val_accuracy: 0.5942 -
Epoch 43/100
504/504 [=====] - ETA: 0s - loss: 1.0081 - accuracy: 0.6395
Epoch 43: val_loss improved from 1.10169 to 1.09355, saving model to /content/drive/MyDrive/DL Facial Recognition /FacialExpression_w
504/504 [=====] - 42s 84ms/step - loss: 1.0081 - accuracy: 0.6395 - val_loss: 1.0935 - val_accuracy: 0.5987 -
Epoch 44/100
504/504 [=====] - ETA: 0s - loss: 0.9997 - accuracy: 0.6444
Epoch 44: val_loss did not improve from 1.09355
504/504 [=====] - 38s 75ms/step - loss: 0.9997 - accuracy: 0.6444 - val_loss: 1.1084 - val_accuracy: 0.6048 -
Epoch 45/100
504/504 [=====] - ETA: 0s - loss: 0.9912 - accuracy: 0.6430
Epoch 45: val_loss improved from 1.09355 to 1.09098, saving model to /content/drive/MyDrive/DL Facial Recognition /FacialExpression_w
504/504 [=====] - 42s 84ms/step - loss: 0.9912 - accuracy: 0.6430 - val_loss: 1.0910 - val_accuracy: 0.6120 -
Epoch 46/100
504/504 [=====] - ETA: 0s - loss: 0.9792 - accuracy: 0.6460
Epoch 46: val_loss did not improve from 1.09098
504/504 [=====] - 38s 74ms/step - loss: 0.9792 - accuracy: 0.6460 - val_loss: 1.1025 - val_accuracy: 0.6148 -
Epoch 47/100
504/504 [=====] - ETA: 0s - loss: 0.9772 - accuracy: 0.6478
Epoch 47: val_loss did not improve from 1.09098
504/504 [=====] - 37s 74ms/step - loss: 0.9772 - accuracy: 0.6478 - val_loss: 1.1394 - val_accuracy: 0.5875 -
Epoch 48/100
504/504 [=====] - ETA: 0s - loss: 0.9735 - accuracy: 0.6511
Epoch 48: val_loss did not improve from 1.09098
504/504 [=====] - 37s 74ms/step - loss: 0.9735 - accuracy: 0.6511 - val_loss: 1.1099 - val_accuracy: 0.6070 -
Epoch 49/100
504/504 [=====] - ETA: 0s - loss: 0.9638 - accuracy: 0.6520
Epoch 49: val_loss did not improve from 1.09098
504/504 [=====] - 37s 74ms/step - loss: 0.9638 - accuracy: 0.6520 - val_loss: 1.0957 - val_accuracy: 0.6031 -
Epoch 50/100
504/504 [=====] - ETA: 0s - loss: 0.9596 - accuracy: 0.6564
Epoch 50: val_loss did not improve from 1.09098
504/504 [=====] - 37s 74ms/step - loss: 0.9596 - accuracy: 0.6564 - val_loss: 1.1052 - val_accuracy: 0.6020 -
Epoch 51/100
504/504 [=====] - ETA: 0s - loss: 0.9547 - accuracy: 0.6557
Epoch 51: val_loss did not improve from 1.09098
504/504 [=====] - 37s 74ms/step - loss: 0.9547 - accuracy: 0.6557 - val_loss: 1.0956 - val_accuracy: 0.5964 -
Epoch 52/100
504/504 [=====] - ETA: 0s - loss: 0.9477 - accuracy: 0.6586
Epoch 52: val_loss improved from 1.09098 to 1.08075, saving model to /content/drive/MyDrive/DL Facial Recognition /FacialExpression_w
504/504 [=====] - 42s 84ms/step - loss: 0.9477 - accuracy: 0.6586 - val_loss: 1.0808 - val_accuracy: 0.6065 -
Epoch 53/100
504/504 [=====] - ETA: 0s - loss: 0.9401 - accuracy: 0.6610
Epoch 53: val loss did not improve from 1.08075
```

```
1 score = model.evaluate(X_test, y_test)
2 print('Test Accuracy: {}'.format(score[1]))
```

```
57/57 [=====] - 2s 18ms/step - loss: 0.9468 - accuracy: 0.6546
Test Accuracy: 0.6545960903167725
```

```
1 accuracy = history.history['accuracy']  
2 val_accuracy = history.history['val_accuracy']  
3 loss = history.history['loss']  
4 val_loss = history.history['val_loss']
```